

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum- 590014, Karnataka.



LAB REPORT

on

Machine Learning (23CS6PCMAL)

Submitted by

Prajwal K K (1BM22CS199)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING

in

COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU - 560019

February 2025 – June 25

B.M.S. College of Engineering

Bull Temple Road, Bangalore 560019

(Affiliated To Visvesvaraya Technological University, Belgaum)

Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Prajwal K K (1BM22CS199)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfilment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Laboratory report has been approved as it satisfies the academic requirements in respect of a Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Lab Faculty Incharge Name: Ms. Saritha A N Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
--	---

INDEX

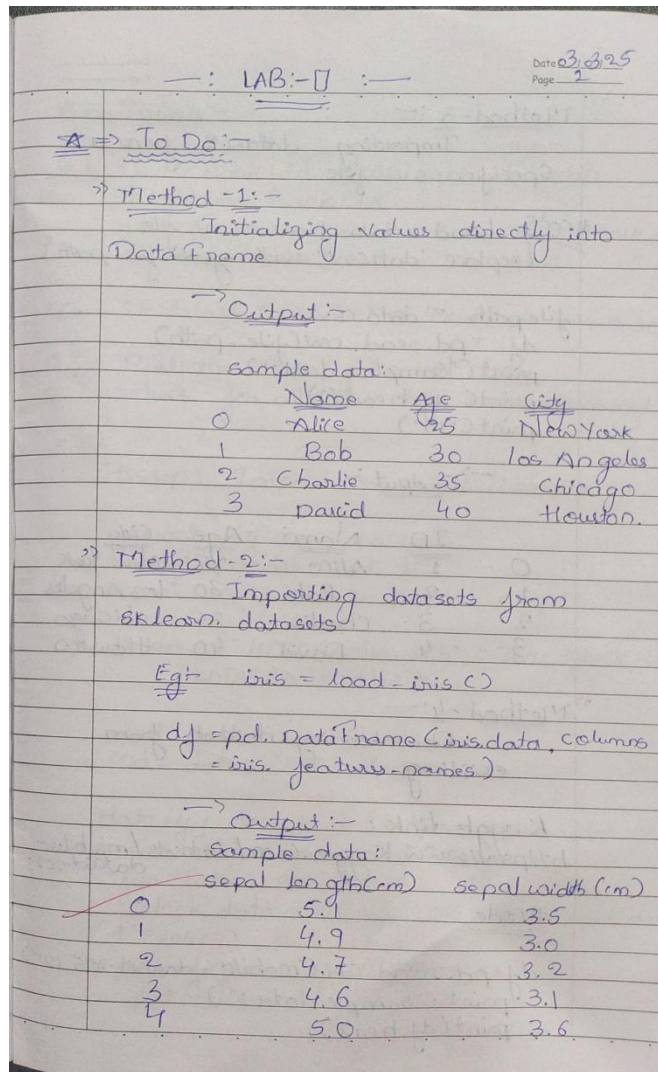
Sl. No.	Date	Experiment Title	Page No.
1	21-02-2025	Write a python program to import and export data using panda's library functions.	1
2	03-03-2025	Demonstrate various data pre-processing techniques for a given dataset.	6
3	10-03-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	11
4	17-03-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset.	20
5	24-03-2025	Build Logistic Regression Model for a given dataset.	24
6	07-04-2025	Build KNN Classification model for a given dataset.	26
7	07-04-2025	Build Support vector machine model for a given dataset.	30
8	21-04-2025	Implement Random forest ensemble method on a given dataset.	34
9	05-05-2025	Implement Boosting ensemble method on a given dataset.	37
10	05-05-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	40
11	12-05-2025	Implement Dimensionality reduction using Principle Component Analysis (PCA) method.	44

Github Link: [PRAJWALKK007/ML-LAB](https://github.com/PRAJWALKK007/ML-LAB)

Program 1

Write a python program to import and export data using Pandas library functions

Screenshot:



Method-3 :-

Importing datasets from a specific csv file

#code Load data from a csv file
[replace 'data.csv' with your file path]

```
file path = 'data.csv'
df = pd.read_csv(file-path)
print("sample data:")
print(df.head(1))
print(df.shape)
```

→ Output :-

	ID	Name	Age	City
0	1	Alice	25	New York
1	2	Bob	30	Los Angeles
2	3	Charlie	35	Chicago
3	4	David	40	Houston

Method-4 :-

Downloading datasets from existing

Kaggle link :

<https://www.kaggle.com/datasets/mobile-dataset-2025>

#code

```
df = pd.read_csv('mobile-dataset-2025.csv')
print("sample data:")
print(df.head())
```

Code:

```
from sklearn.datasets import load_iris

import pandas as pd

iris = load_iris()

df = pd.DataFrame(iris.data, columns=iris.feature_names)

df.head()


df['target'] = iris.target

df

import kagglehub

# Download latest version

path = kagglehub.dataset_download("abdulmalik1518/mobiles-dataset-2025")

print("Path to dataset files:", path)


df = pd.read_csv("/content/Mobiles_Dataset_(2025).csv", encoding='latin-1') # or 'ISO-8859-1', or
'cp1252'

df.head()

df['Company Name']


data = {"USN" : ['1', "2", "3"], "Name" : ["A", "B", "C"]}

df = pd.DataFrame(data)

df
```

```
from sklearn.datasets import load_diabetes

diabetes = load_diabetes()

df = pd.DataFrame(diabetes.data, columns=diabetes.feature_names)

df.head()

df.columns

df = pd.read_csv("/content/Dataset_of_Diabetes .csv")

df.head()

import yfinance as yf
import pandas as pd

import matplotlib.pyplot as plt
```

```
tickers = ["RELIANCE.NS", "TCS.NS", "INFY.NS"]

# Fetch historical data for the last 1 year

data = yf.download(tickers, start="2022-10-01", end="2023-10-01", group_by='ticker')

# Display the first 5 rows of the dataset

print("First 5 rows of the dataset:")

print(data.head())

print("\nShape of the dataset:")

print(data.shape)
```

```
# Summary statistics for a specific stock (e.g., Reliance)

reliance_data = data['RELIANCE.NS']

print("\nSummary statistics for Reliance Industries:")

print(reliance_data.describe())

# Calculate daily returns

reliance_data['Daily Return'] = reliance_data['Close'].pct_change()


# Plot the closing price and daily returns

plt.figure(figsize=(12, 6))

plt.subplot(2, 1, 1)

reliance_data['Close'].plot(title="Reliance Industries - Closing Price")

plt.subplot(2, 1, 2)

reliance_data['Daily Return'].plot(title="Reliance Industries - Daily Returns", color='orange')

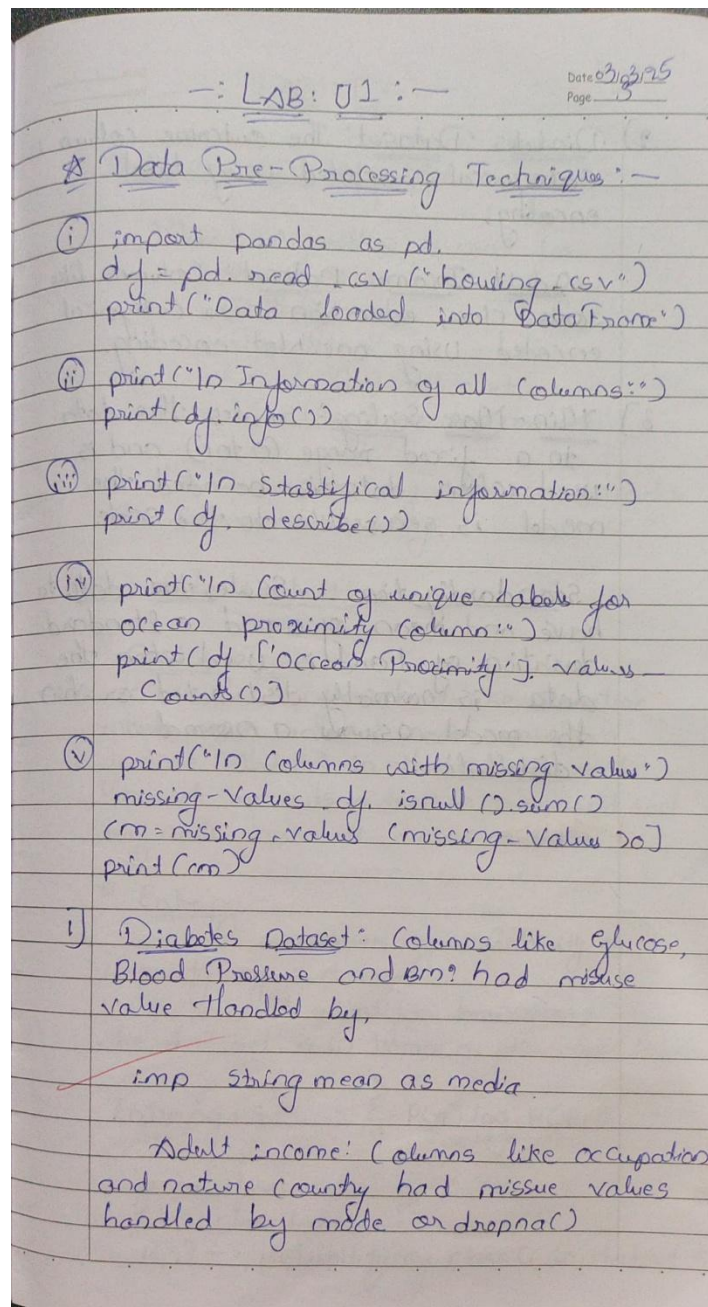
plt.tight_layout()

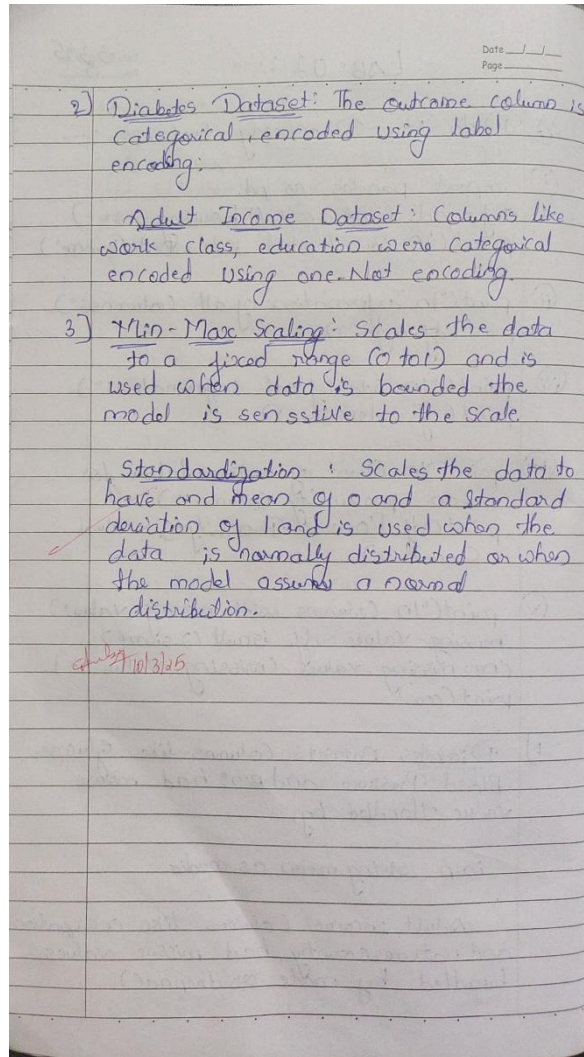
plt.show()
```


Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot:





Code:

```
import pandas as pd

import numpy as np

# Load dataset

df = pd.read_csv("data.csv")

print(df.head())
```

```
# Check missing values
```

```
print(df.isnull().sum())
```

```
# Drop rows with missing values
```

```
df_cleaned = df.dropna()
```

```
# Or fill missing values with mean/median
```

```
df['Age'].fillna(df['Age'].mean(), inplace=True)
```

```
df['Salary'].fillna(df['Salary'].median(), inplace=True)
```

```
# For nominal categories
```

```
df = pd.get_dummies(df, columns=['Gender', 'Country'], drop_first=True)
```

```
# For ordinal categories
```

```
from sklearn.preprocessing import OrdinalEncoder
```

```
encoder = OrdinalEncoder()
```

```
df[['Education_Level']] = encoder.fit_transform(df[['Education_Level']])
```

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```
# Standardization (Z-score)
```

```
scaler = StandardScaler()
```

```
df[['Age', 'Salary']] = scaler.fit_transform(df[['Age', 'Salary']])
```

```
# Min-Max Normalization
```

```
minmax = MinMaxScaler()
```

```
df[['Age', 'Salary']] = minmax.fit_transform(df[['Age', 'Salary']])
```

```
# Using IQR method
```

```
Q1 = df['Salary'].quantile(0.25)
```

```
Q3 = df['Salary'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
df = df[(df['Salary'] >= Q1 - 1.5*IQR) & (df['Salary'] <= Q3 + 1.5*IQR)]
```

```
df['Age_Salary_Ratio'] = df['Age'] / df['Salary']
```

```
# Drop irrelevant columns
```

```
df.drop(['User_ID', 'Name'], axis=1, inplace=True)
```

```
# Correlation-based filtering
```

```
correlation_matrix = df.corr()
```

```
print(correlation_matrix)
```

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop('Purchased', axis=1)
```

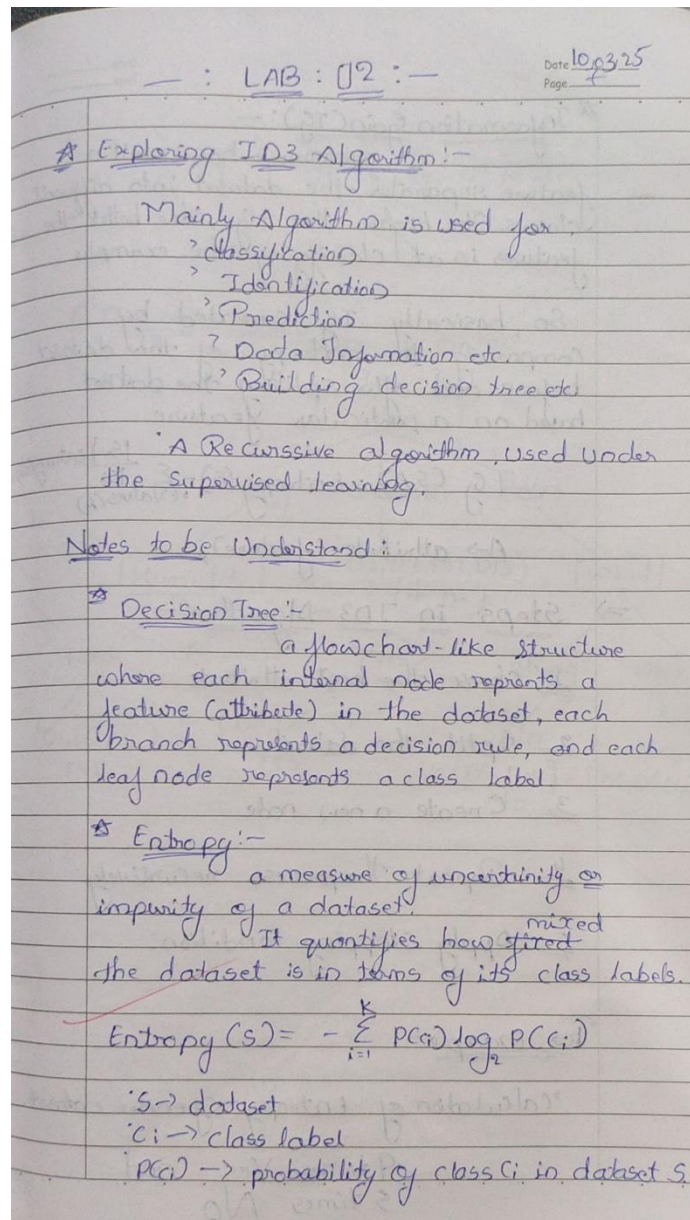
```
y = df['Purchased']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Program 3

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample

Screenshot:



Date: / /
Page:
* Information Gain (IG): -

measures how well a feature separates the dataset into different classes. The higher the IG, the better the feature is at classifying the examples.

So, basically IG created by comparing the entropy of the dataset before and after splitting the dataset based on a particular feature.

$$IG(S, A) = Entropy(S) - \sum_{v \in \text{values}(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

A \rightarrow attribute (-feature),

\Rightarrow Steps in ID3 Algorithm:-

1. Choose the best attribute

2. Split the dataset

3. Create a new node

4. Repeat the process recursively

5. Apply stopping condition.

For Example:-

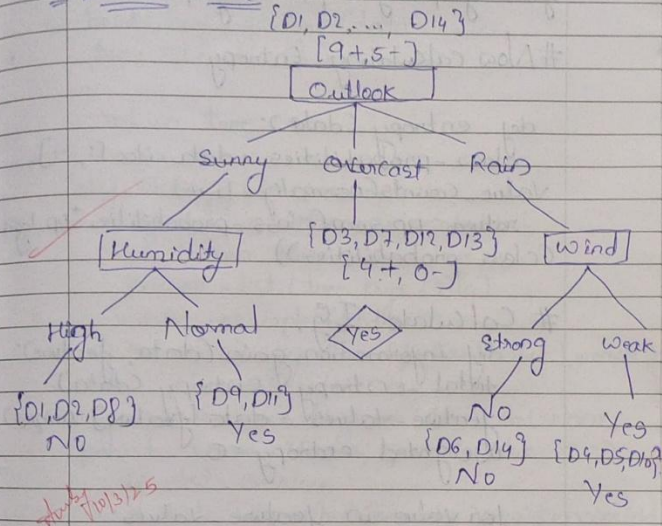
Calculation of Entropy of entire dataset.

9 times Yes and
5 times No

$$(S) = - \left(\frac{9}{14} \log_2 \frac{9}{14} + \frac{5}{14} \log_2 \frac{5}{14} \right)$$

$$S \approx 0.940$$

Decision Tree



17/03/20

ID3 code:

```
import pandas as pd
import numpy as np

def entropy(data):
    class_prob = data.iloc[:, -1].value_counts(normalize=True)
    return -np.sum(class_prob * np.log2(class_prob))

def information_gain(data, features):
    total_entropy = entropy(data)
    feature_values = data[features].unique()
    weighted_entropy = 0
    for value in feature_values:
        subset = data[data[features] == value]
        weighted_entropy += (len(subset) / len(data)) * entropy(subset)
    return total_entropy - weighted_entropy

def best_feature(data):
    features = data.columns[:-1]
    gains = {feature: information_gain(data, feature) for feature in features}
    return max(gains, key=gains.get)

def id3(data, feature1=None):
    if len(data.iloc[:, -1].unique()) == 1:
        return data.iloc[:, -1].iloc[0]

    if len(features) == 0:
```

* Implementing the ID3 Algorithm :-

#Code

```
import pandas as pd
import numpy as np
from graphviz import Digraph

# Now calculating Entropy

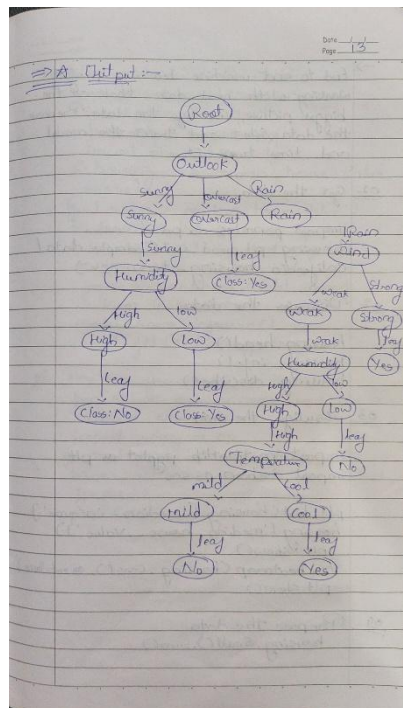
def entropy(data):
    class_probabilities = data.iloc[:, -1].value_counts(normalize=True)
    return -np.sum(class_probabilities * np.log2(class_probabilities))

# Calculate IG
def information_gain(data, feature):
    total_entropy = entropy(data)
    feature_values = data[feature].unique()
    weighted_entropy = 0

    for value in feature_values:
        subset = data[data[feature] == value]
        weighted_entropy += (len(subset) / len(data)) * entropy(subset)

    return total_entropy - weighted_entropy

def best_feature(data):
    features = data.columns[:-1]
    gains = {feature: information_gain(data, feature) for feature in features}
    return max(gains, key=gains.get)
```



Code:

```
import pandas as pd

import numpy as np

from graphviz import Digraph


# Calculate Entropy

def entropy(data):

    class_probabilities = data.iloc[:, -1].value_counts(normalize=True)

    return -np.sum(class_probabilities * np.log2(class_probabilities))


# Calculate Information Gain

def information_gain(data, feature):

    total_entropy = entropy(data)

    feature_values = data[feature].unique()

    weighted_entropy = 0

    for value in feature_values:

        subset = data[data[feature] == value]

        weighted_entropy += (len(subset) / len(data)) * entropy(subset)

    return total_entropy - weighted_entropy


# Find the best feature to split the data

def best_feature(data):

    features = data.columns[:-1] # Exclude the target column

    gains = {feature: information_gain(data, feature) for feature in features}
```

```

return max(gains, key=gains.get)

# Create the decision tree

def id3(data, features=None):

    if len(data.iloc[:, -1].unique()) == 1: # All data points belong to the same class

        return data.iloc[:, -1].iloc[0]

    if len(features) == 0: # No more features to split on

        return data.iloc[:, -1].mode()[0]

    best = best_feature(data)

    tree = {best: {}}

    new_features = features.copy()

    new_features.remove(best)

    for value in data[best].unique():

        subset = data[data[best] == value]

        tree[best][value] = id3(subset, new_features)

    return tree

# Function to classify new examples based on the decision tree

def classify(tree, example):

```

```

if not isinstance(tree, dict):

    return tree

feature = list(tree.keys())[0]

value = example[feature]

return classify(tree[feature][value], example)

```

Function to visualize the decision tree using Graphviz

```
def create_tree_diagram(tree, dot=None, parent_name="Root", parent_value=""):
```

```
    if dot is None:
```

```
        dot = Digraph(format="png", engine="dot")
```

```
    if isinstance(tree, dict): # Tree node
```

```
        for feature, branches in tree.items():
```

```
            feature_name = f"{parent_name}_{feature}"
```

```
            dot.node(feature_name, feature)
```

```
            dot.edge(parent_name, feature_name, label=parent_value)
```

```
        for value, subtree in branches.items():
```

```
            value_name = f"{feature_name}_{value}"
```

```
            dot.node(value_name, f"{feature}: {value}")
```

```
            dot.edge(feature_name, value_name, label=str(value))
```

```
        # Recurse for each subtree
```

```
        create_tree_diagram(subtree, dot, value_name, str(value))
```

```
    else: # Leaf node
```

```

dot.node(parent_name + "_class", f"Class: {tree}")

dot.ede(parent_name, parent_name + "_class", label="Leaf")

return dot

# Example usage

data = pd.DataFrame({

    'Outlook': ['Sunny', 'Sunny', 'Overcast', 'Rain', 'Rain', 'Rain', 'Overcast', 'Sunny', 'Sunny', 'Rain',
'Sunny', 'Overcast', 'Overcast', 'Rain'],

    'Temperature': ['Hot', 'Hot', 'Hot', 'Mild', 'Cool', 'Cool', 'Cool', 'Mild', 'Cool', 'Mild', 'Mild', 'Mild', 'Hot',
'Mild'],

    'Humidity': ['High', 'High', 'High', 'High', 'High', 'Low', 'Low', 'High', 'Low', 'Low', 'Low', 'High', 'Low',
'High'],

    'Wind': ['Weak', 'Strong', 'Weak', 'Weak', 'Weak', 'Weak', 'Strong', 'Weak', 'Weak', 'Strong', 'Strong',
'Weak', 'Strong', 'Weak'],

    'PlayTennis': ['No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'Yes', 'No', 'Yes', 'Yes', 'Yes', 'Yes', 'Yes', 'No']

})

# Train the decision tree

tree = id3(data, features=list(data.columns[:-1]))

print("Decision Tree:", tree)

# Classify a new example

example = {'Outlook': 'Sunny', 'Temperature': 'Cool', 'Humidity': 'Low', 'Wind': 'Strong'}

prediction = classify(tree, example)

print("Prediction for the example:", prediction)

# Visualize the decision tree

dot = create_tree_diagram(tree)

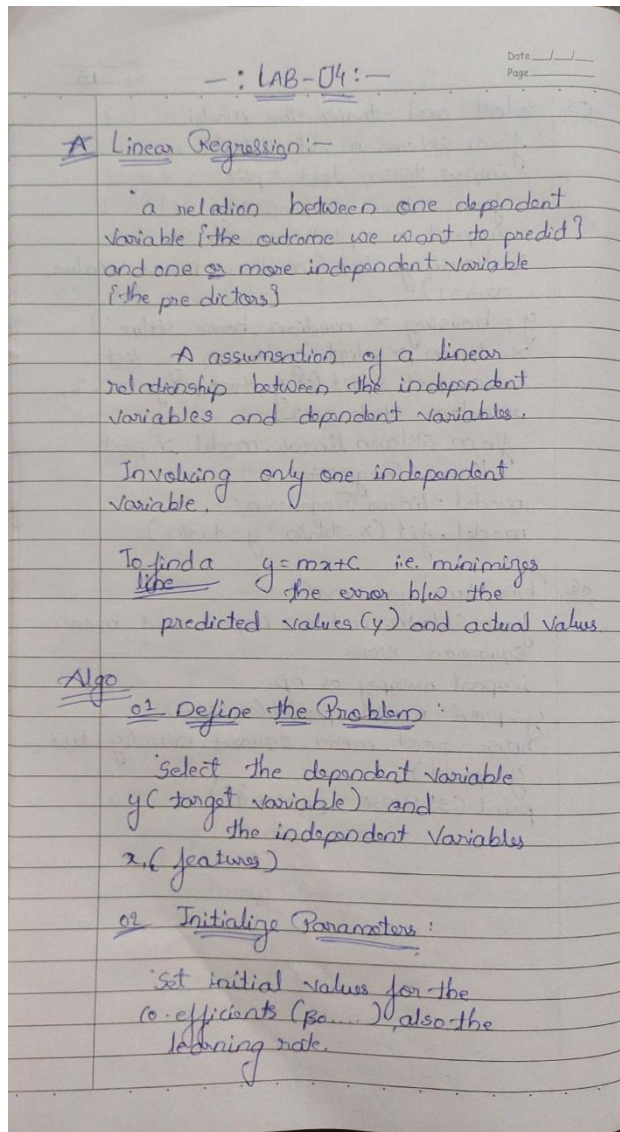
dot.render("decision_tree", view=True) # This will generate and open the tree diagram

```

Program 4

Implement Linear and Multi-Linear Regression algorithm for appropriate dataset

Screenshot:



MSE \rightarrow Average of squared difference b/w actual and predicted values.

Date: / /
Page: 15

03. Compute Prediction:

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \epsilon$$

04. Calculate MSE: {Mean Squared Error}

$$MSE = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

$N \rightarrow$ observations

$\hat{y}_i \rightarrow$ predicted values

$y_i \rightarrow$ actual values

* Applications:

- Analyzing risk in financial systems
- Forecasting sales or revenue
- Estimating trends in data
- Predicting Student Satisfaction

* Pseudocode for Linear Regression:

Function LinearRegression(x, y):

#step 1: Add a column of ones to x for the intercept term

x = AddColumnOfOnes(x)

#step 2: Compute the coefficients using the OLS formula

#beta = $(x^T \cdot x)^{-1} \cdot x^T \cdot y$

x-transpose = Transpose(x)

xTx = Multiply(x-transpose, x)

xTx.inverse = Inverse(xTx)

Code:

Linear Regression

```
import pandas as pd

df = pd.read_csv("/content/tvmarketing.csv")

df

# Visualise the relationship between the features and the response using scatterplots

df.plot(x='TV',y='Sales',kind='scatter')


from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(df['TV'], df['Sales'], test_size=0.2, random_state=42)

from sklearn.linear_model import LinearRegression model = LinearRegression()
model.fit(x_train.values.reshape(-1, 1), y_train) y_train
model.coef_

model.intercept_
```

MultiLinearRegression

```
import pandas as pd

# Step 2 : import data

house = pd.read_csv('https://github.com/YBIFoundation/Dataset/raw/main/Boston.csv')

# display first 5 rows
```

```
house.head()
```

```
y = house['MEDV']
```

```
X = house.drop(['MEDV'],axis=1)
```

```
# Step 4 : train test split
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.7, random_state=2529)
```

```
# Step 5 : select model
```

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
# Step 6 : train or fit model
```

```
model.fit(X_train,y_train)
```

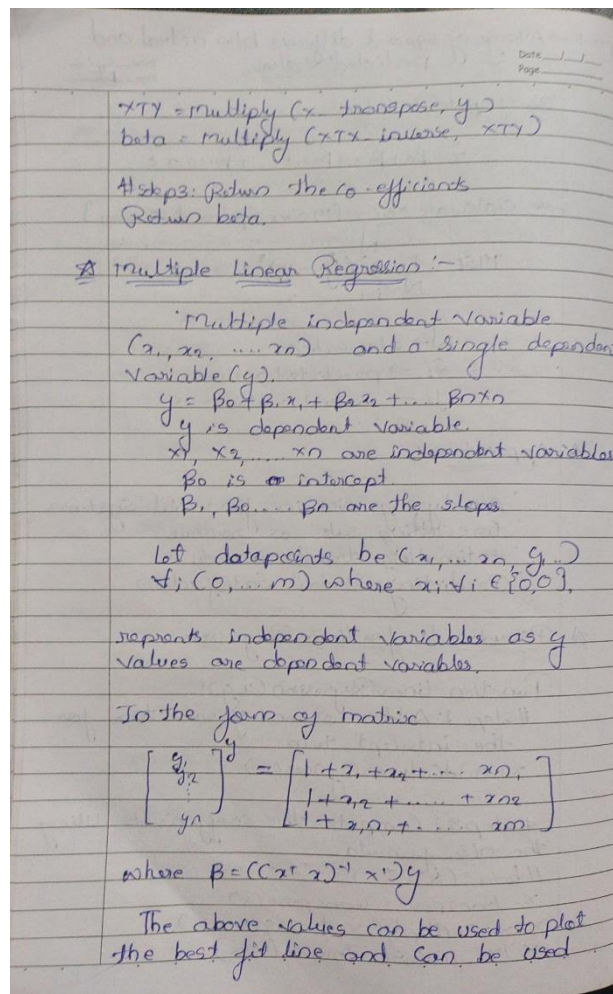
```
model.intercept_
```

```
model.coef_
```

Program 5

Build Logistic Regression Model for a given dataset

Screenshot:



to predict future values.

★ Logistic Regression :-

- Logistic regression approach operates on Sigmoid curve rather than best fit line, we get a value $\in [0, 1]$ (Binary Classification) and then classify into +ve or -ve by comparing with median.

Let data points be $(x_i, y_i) \forall i \in [0, n]$,
finding best fit line through
previously mentioned methods

$$V = \frac{1}{1 + e^{-(mx+c)}} \times \frac{1}{1 + e^{-(b_1m + b_2x)}}$$

Classification will be based on - the
obtained value V.

- * If $V < 0.5 \rightarrow$ then "no".
- * If $V > 0.5 \rightarrow$ then "yes".

Code:

```
from sklearn.linear_model import LogisticRegression

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.metrics import accuracy_score


# Load sample dataset (binary classification - Iris with only 2 classes)

iris = load_iris()

X = iris.data[iris.target != 2]

y = iris.target[iris.target != 2]


# Train/Test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)


# Logistic Regression model

model = LogisticRegression()

model.fit(X_train, y_train)


# Predict and evaluate

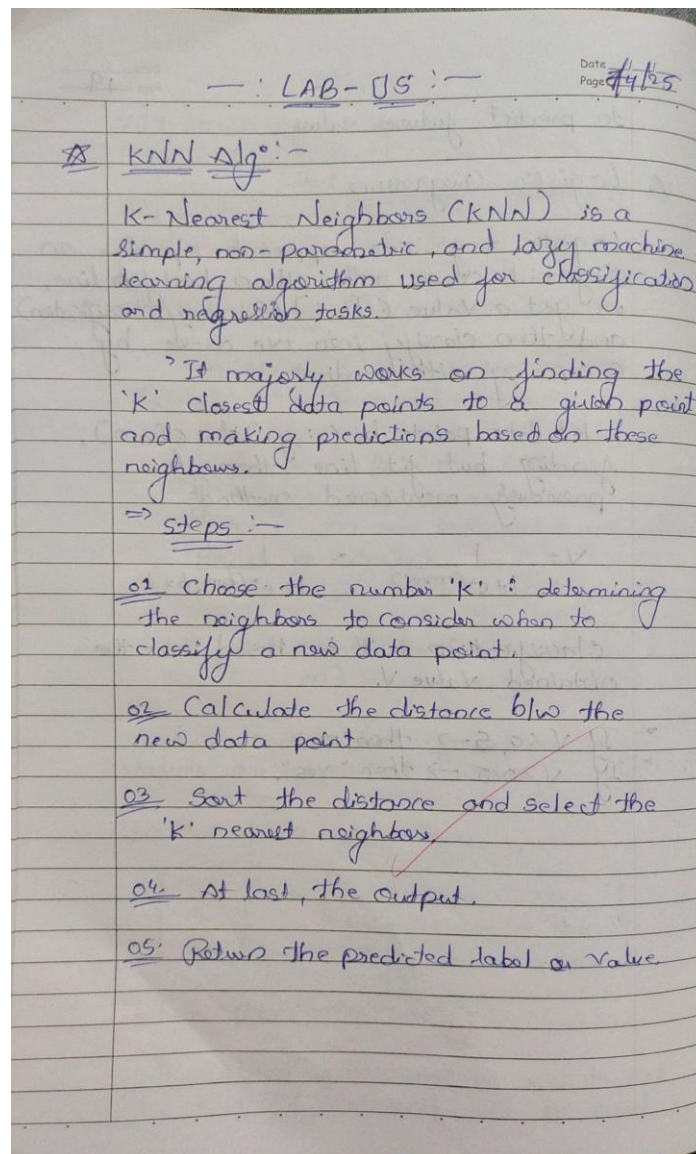
y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
```

Program 6

Build KNN Classification model for a given dataset

Screenshot:



$$d = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2} \text{ \{Distance Metric\}}$$

Date: 1/1
Page: 21

Code Using sklearn

iris = load_iris()

x = iris.data

y = iris.target

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

scaler = StandardScaler()

x_train_scaled = scaler.fit_transform(x_train)

x_test_scaled = scaler.transform(x_test)

knn = KNeighborsClassifier(n_neighbors=3)

knn.fit(x_train_scaled, y_train)

y_pred = knn.predict(x_test_scaled)

accuracy = accuracy_score(y_test, y_pred)

=> Output:-

Accuracy of KNN Classifier: 1.00

Predictions: [1 0 2 1 0 1 2 0 0 0

1 2 1 2 0 2 2 2 2 0]

True labels: [1 0 2 1 0 1 2 1 2 0 0 0

1 2 1 2 0 2 2 2 2 0]

Tuning 'K':

If K is too small, the model may be noisy and overfit the data (high variance)

Code:

KNN

```
import numpy as np

from collections import Counter

class KNN:

    def __init__(self, k=3): self.k = k

    def fit(self, X, y):

        self.X_train = np.array(X)

        self.y_train = np.array(y)

    def euclidean_distance(self, x1, x2):

        return np.sqrt(np.sum((x1 - x2) ** 2))

    def predict(self, X):

        predictions = [self._predict(x) for x in X]

        return np.array(predictions)

    def _predict(self, x):

        # Compute distances to all training points

        distances = [self.euclidean_distance(x, x_train) for x_train in self.X_train]

        # Get indices of k nearest neighbors
```



```

k_indices = np.argsort(distances)[:self.k]

# Get the labels of those neighbors
k_nearest_labels = [self.y_train[i] for i in k_indices]

# Return the most common label
most_common = Counter(k_nearest_labels).most_common(1)

return most_common[0][0]

# Sample dataset (like a mini version of Iris)
X_train = [[1, 2], [2, 3], [3, 1], [6, 5], [7, 7], [8, 6]]
y_train = [0, 0, 0, 1, 1, 1]

# Test data
X_test = [[5, 5], [1, 1]]

# Using the KNN modelh
knn = KNN(k=3)
knn.fit(X_train, y_train)
predictions = knn.predict(X_test)

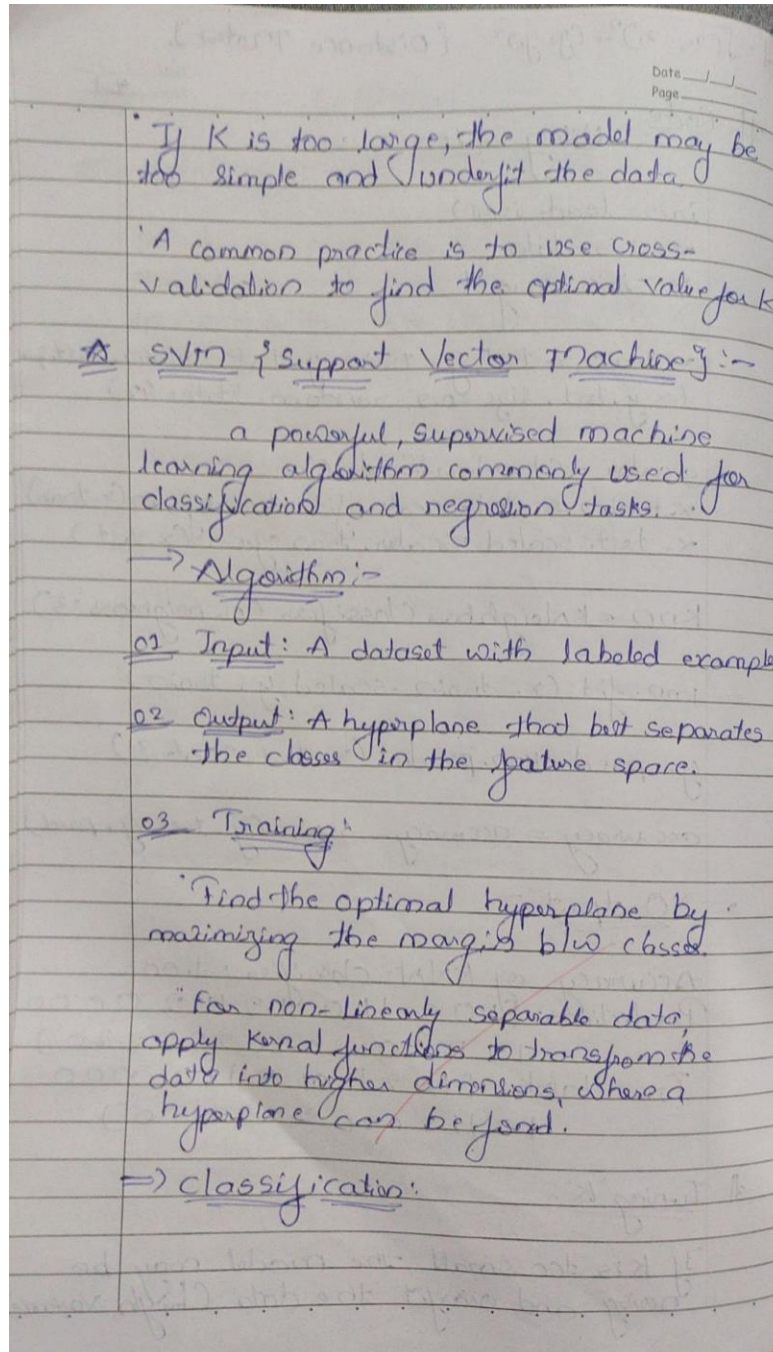
print("Predictions:", predictions)

```

Program 7

Build Support vector machine model for a given dataset

Screenshot:



1 → Yes
0 → No

Date ____/____/____
Page ____

Input: A set of labeled data points.

Output: A class label based on the optimal hyperplane.

Goal: Maximize the margin between classes while minimizing misclassification.

#Code:-

For Example:

• Age (in years)

• Income (in thousands of dollars)

• Product Usage Frequency (scale from 1 to 10)

```
np.random.seed(42)
```

```
n_samples = 1000
```

```
age = np.random.randint(18, 70, n_samples)
```

```
income = np.random.randint(30, 150, n_samples)
```

```
usage_freq = np.random.randint(1, 11, n_samples)
```

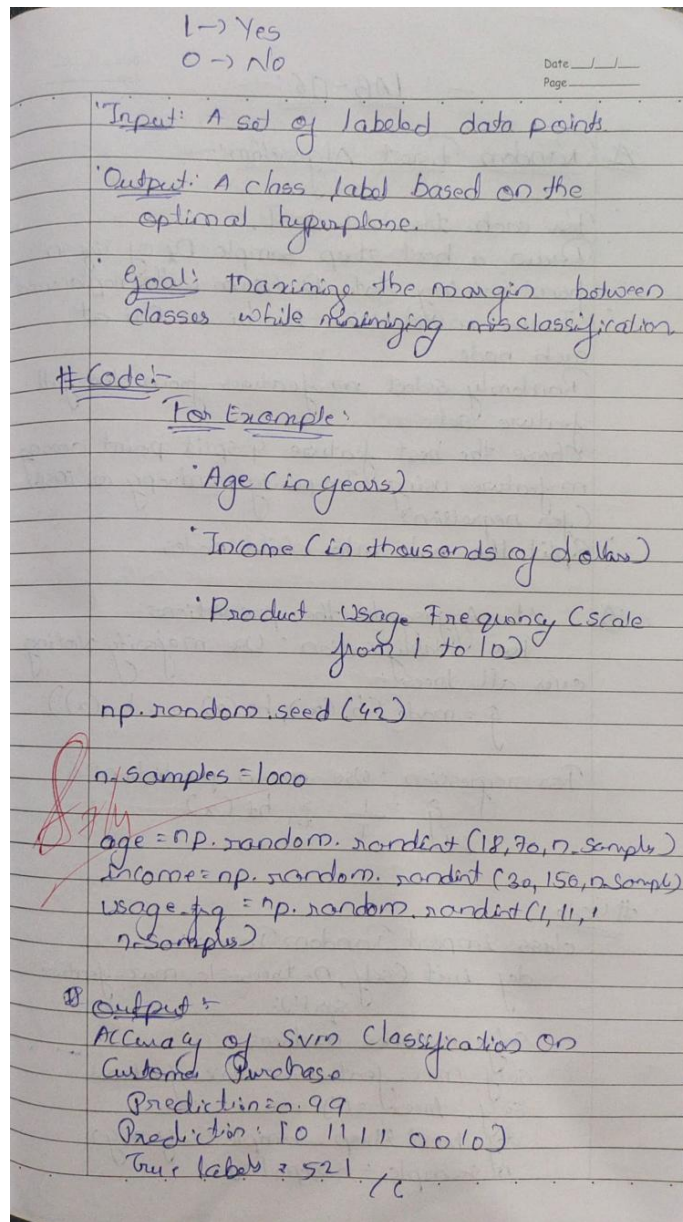
Output:-

Accuracy of SVM Classification on Customer Purchase

Prediction: 0.99

Prediction: [0 1 1 1 0 0 1 0]

True labels: 521/521



Code:

```
from sklearn import datasets

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

import matplotlib.pyplot as plt
```

```
from sklearn.decomposition import PCA

# Load dataset

iris = datasets.load_iris()

X = iris.data

y = iris.target

# For binary classification (class 0 vs 1)

X = X[y != 2]

y = y[y != 2]

# Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Train SVM

clf = SVC(kernel='linear') # Try 'rbf', 'poly', etc.

clf.fit(X_train, y_train)

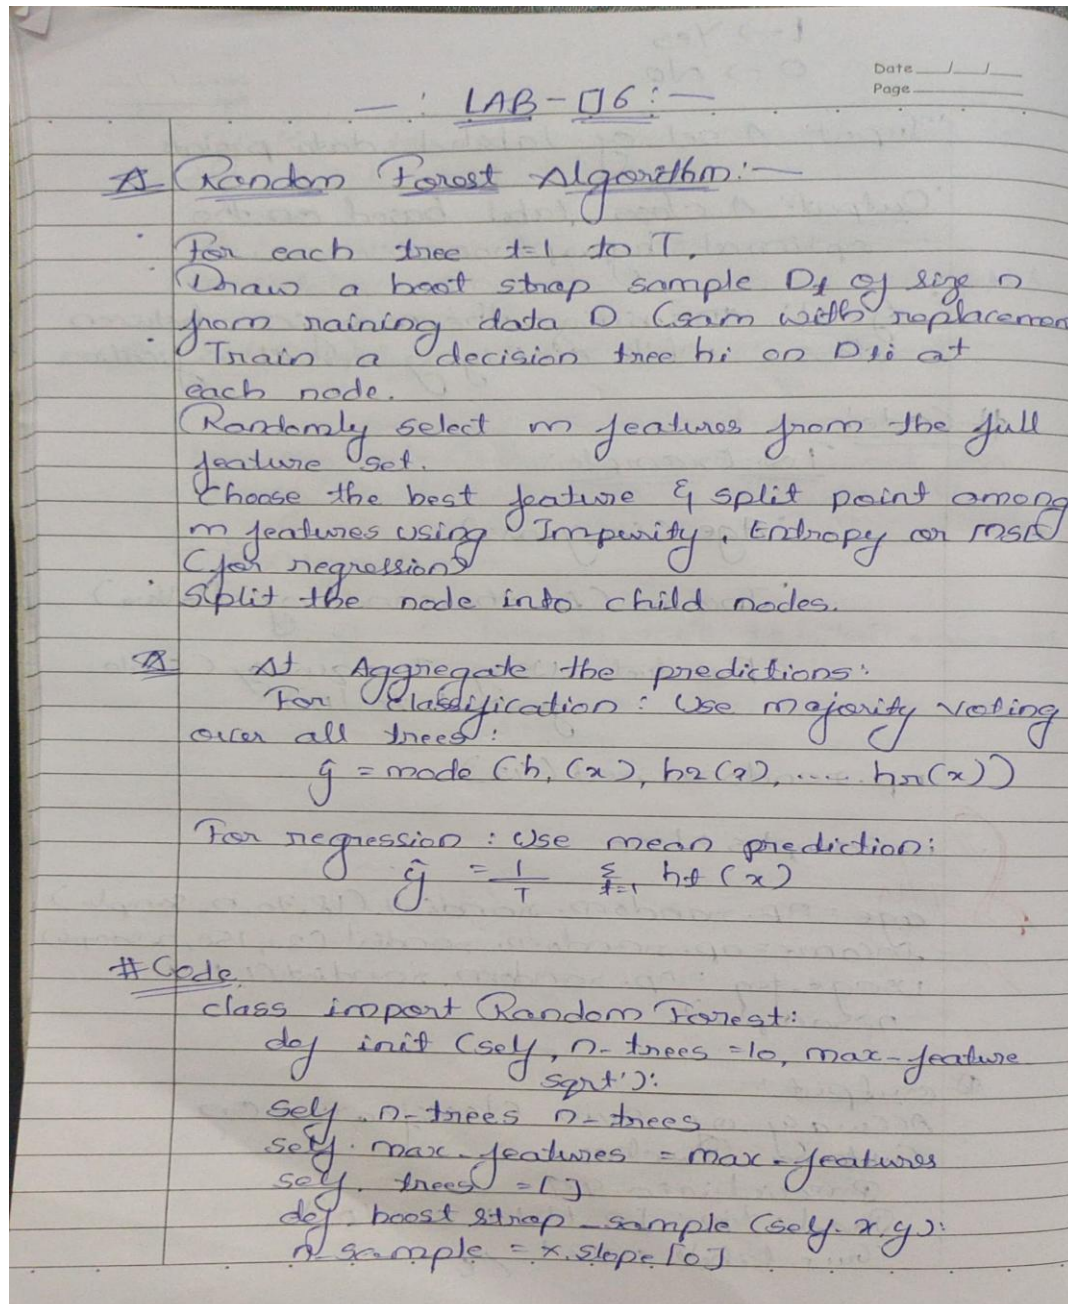
# Accuracy

print("Test Accuracy:", clf.score(X_test, y_test))
```


Program 8

Implement Random forest ensemble method on a given dataset

Screenshot:



```

size = n-sample, replace = True
return x[indices], y[indices]
def get_max_features(self, n_features):
    if self.max_features == 'sqrt':
        return int(np.sqrt(n_features))
    elif is_instances(self, max_features, int):
        return self.max_features
def fit(self, X, y):
    self.tree = 1
    n_features = X.shape[1]
    max_feature = self.get_max_features

```

★ ⇒ K-Means Algorithm:-

01 Initialize centroids:

Randomly choose K data points from x as initial cluster centroids:

$$\mu_1, \mu_2, \dots, \mu_K$$

02 Repeat until convergence.

① Repeat each data point to the nearest centroid. For each point x_i , find the closest centroid μ_j based on distance

$$\mu_j = \frac{1}{|C_i|} \sum_{x_i \in C_i} x_i$$

03 check for convergence

- * if cluster assignments don't change
- * (Centroids don't move significantly, the stop

Code:

```
from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score

# Load sample dataset

iris = load_iris()

X, y = iris.data, iris.target


# Train/test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Initialize Random Forest

rf = RandomForestClassifier(n_estimators=100, random_state=42)

rf.fit(X_train, y_train)


# Predict and evaluate

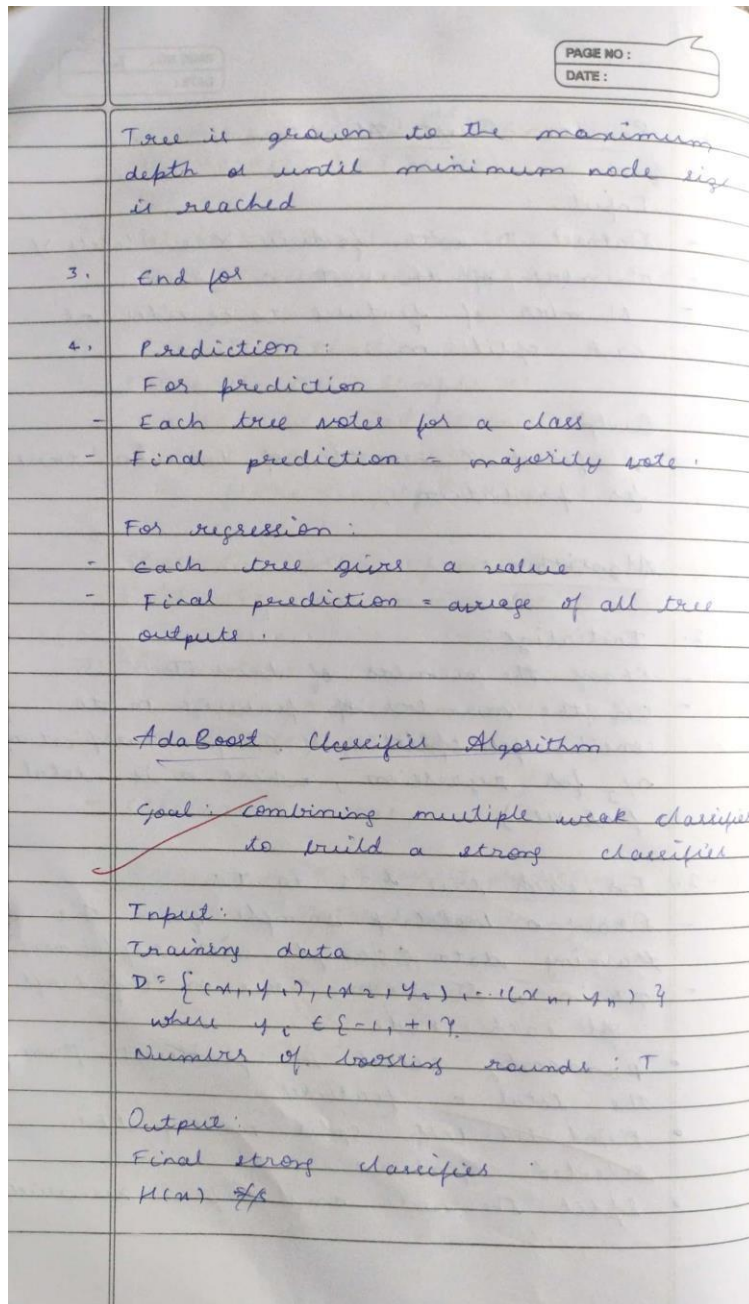
y_pred = rf.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
```


Program 9

Implement Boosting ensemble method on a given dataset

Screenshot:



Code:

```
from sklearn.ensemble import AdaBoostClassifier

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.metrics import accuracy_score


# Load Iris dataset

iris = load_iris()

X, y= iris.data, iris.target
```

```
# For AdaBoost, we'll use binary classification #
```

```
Convert to binary (setosa vs. not-setosa)
```

```
y = (y == 0).astype(int)
```

```
# Split data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train AdaBoost
```

```
model = AdaBoostClassifier(n_estimators=50, learning_rate=1.0, random_state=42)
```

```
model.fit(X_train, y_train)
```

```
# Predict and evaluate
```

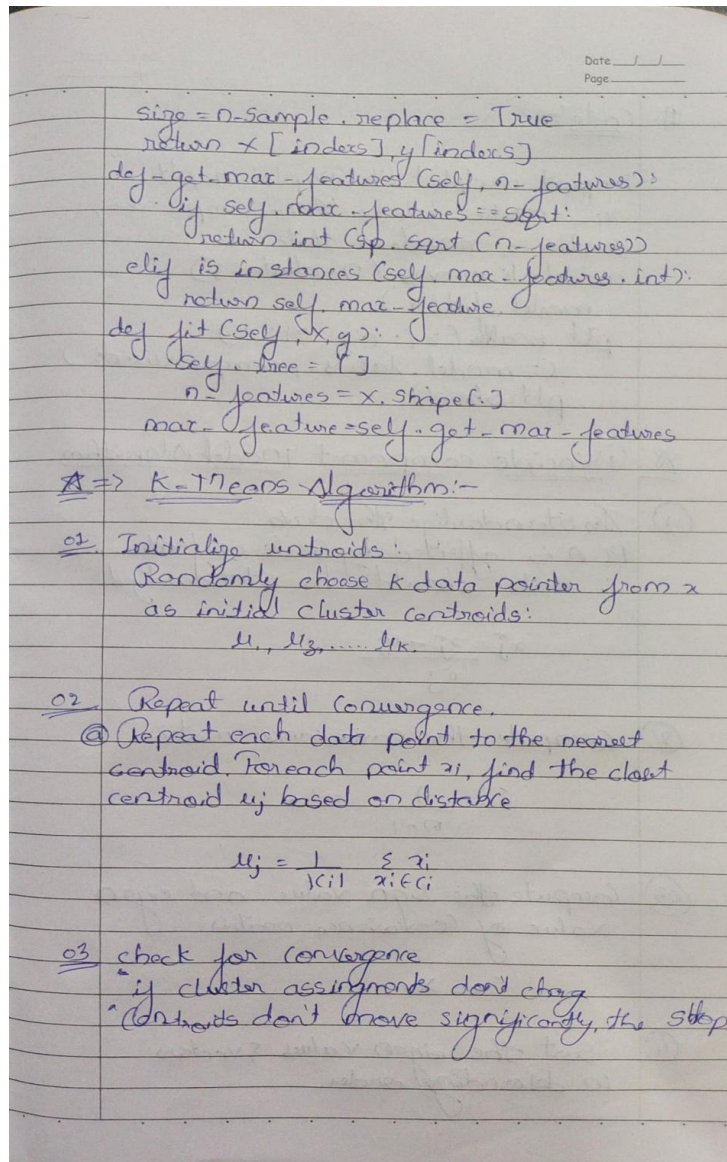
```
y_pred = model.predict(X_test)
```

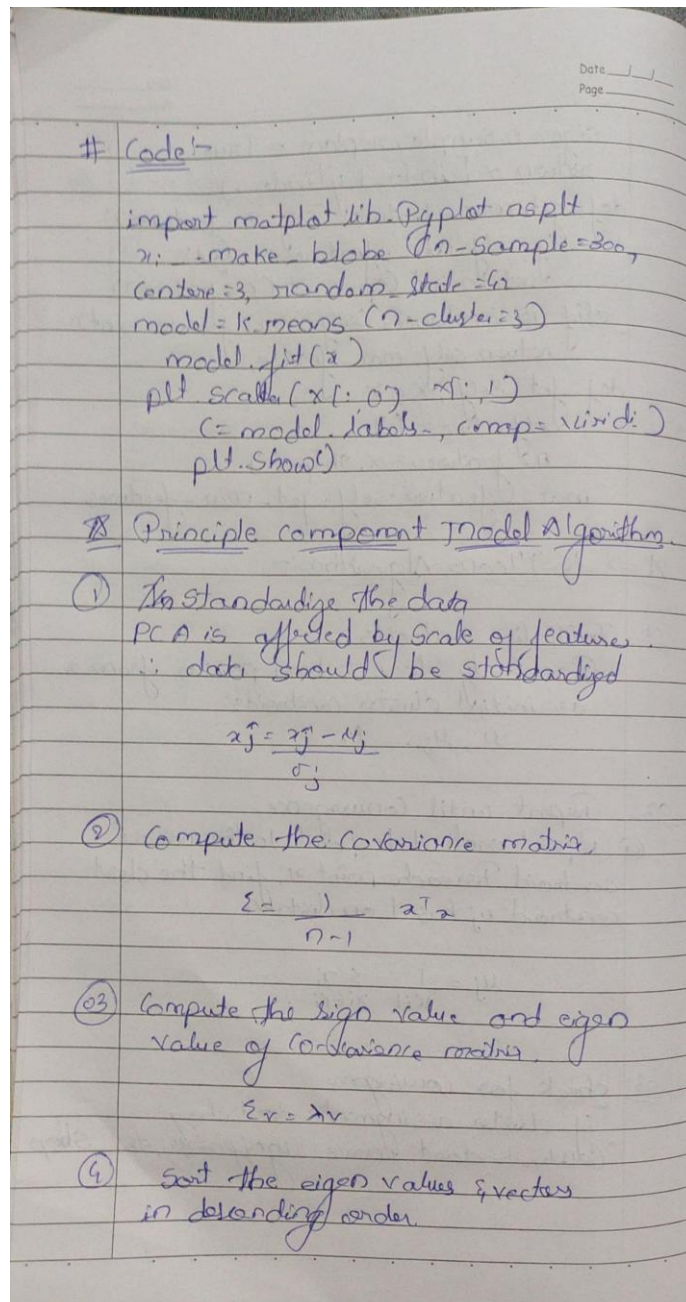
```
print("AdaBoost Accuracy (sklearn):", accuracy_score(y_test, y_pred))
```

Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file

Screenshot:





Code:

import pandas as pd

from sklearn.cluster import KMeans

```

import matplotlib.pyplot as plt

from sklearn.datasets import load_iris # Import load_iris


# Step 1: Load the Iris dataset directly

iris = load_iris()

# Create a DataFrame from the data and target

data = pd.DataFrame(data=iris.data, columns=iris.feature_names)

# Add the target column for potential reference, though not used for clustering

data['target'] = iris.target


# Step 2: Extract only numeric columns (or select required features)

# All features in the Iris dataset are numeric

X = data[iris.feature_names].values # Use the feature names to select columns


# Step 3: Apply KMeans

# Adjust n_clusters based on the expected number of clusters in your data (3 for Iris)

kmeans = KMeans(n_clusters=3, random_state=42, n_init=10) # Added n_init to suppress future
warnings

data['Cluster'] = kmeans.fit_predict(X)


# Step 4: Plot clusters (for 2D data)

# Iris data has 4 features. We will plot the first two features for visualization.

if X.shape[1] >= 2:

```

```
plt.scatter(X[:, 0], X[:, 1], c=data['Cluster'], cmap='viridis')

plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], color='red', marker='x', s=200)

plt.title("K-Means Clustering of Iris Dataset")

plt.xlabel(iris.feature_names[0]) # Label with actual feature name

plt.ylabel(iris.feature_names[1]) # Label with actual feature name

plt.show()

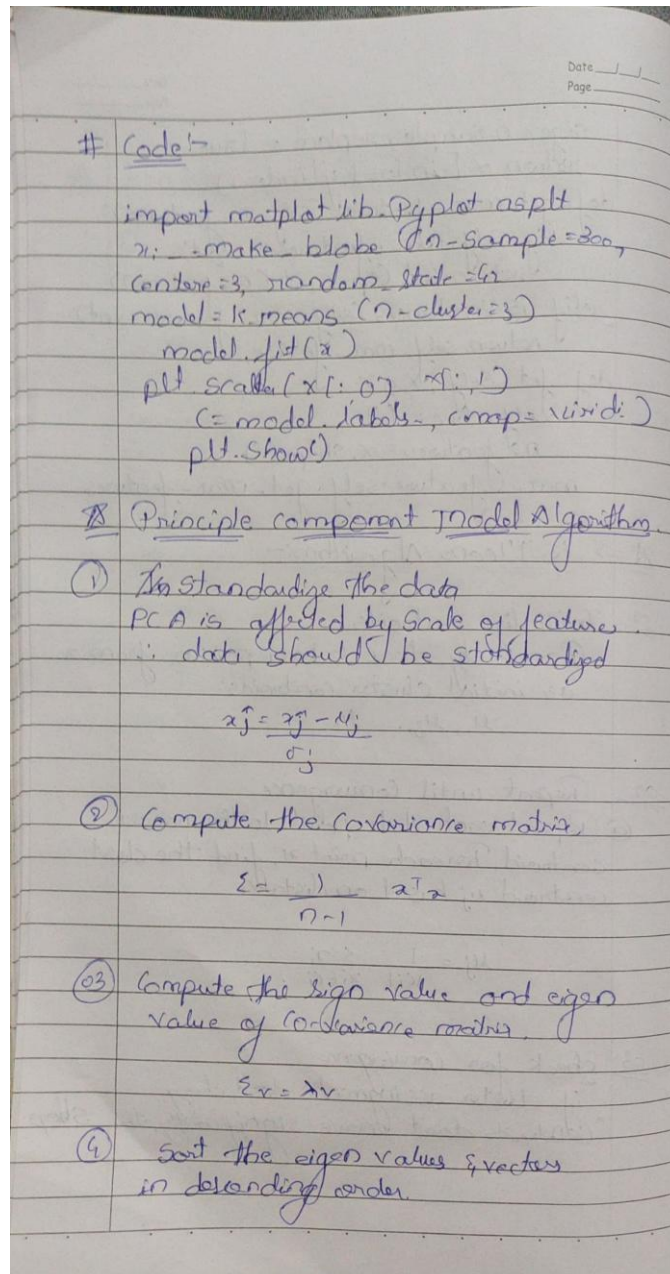
else:

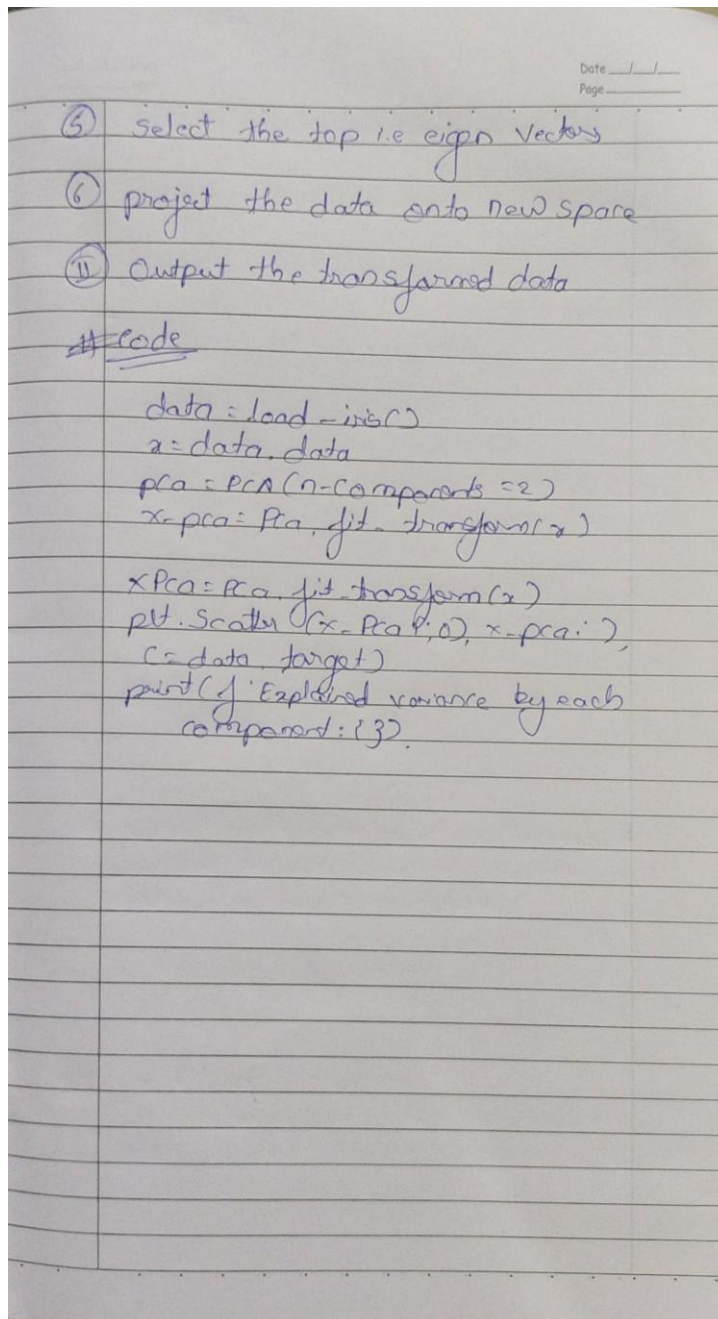
    print("Cannot plot clustering results directly for data with less than 2 features.")
```

Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA) method

Screenshot:





Code:

```
import pandas as pd

from sklearn.decomposition import PCA

from sklearn.preprocessing import StandardScaler
```

```
import matplotlib.pyplot as plt

# Load dataset

data = pd.read_csv("your_data.csv") # Replace with your file

X = data.select_dtypes(include=['float64', 'int64'])

# Step 1: Standardize

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)

# Step 2: Apply PCA

pca = PCA(n_components=2)

X_pca = pca.fit_transform(X_scaled)

# Print explained variance ratio

print("Explained variance ratio:", pca.explained_variance_ratio_)

# Visualize

plt.scatter(X_pca[:, 0], X_pca[:, 1], c='blue', alpha=0.5)

plt.title("PCA - 2D Projection")

plt.xlabel("Principal Component 1")

plt.ylabel("Principal Component 2")

plt.show()
```



Accuracy Before PCA:

Logistic Regression: 0.9016

SVM: 0.8525

Random Forest: 0.8361



Accuracy After PCA (n_components=5):

Logistic Regression: 0.8689

SVM: 0.8689

Random Forest: 0.8852